
Chapter 8: Eigenvalues and Singular Values

Uri M. Ascher and Chen Greif
Department of Computer Science
The University of British Columbia
{ascher,greif}@cs.ubc.ca

Slides for the book

A First Course in Numerical Methods (published by SIAM, 2011)
<http://www.ec-securehost.com/SIAM/CS07.html>

Some slides are from the course of Dr. Peter Arbenz, ETH.

Goals of this chapter

- To find out how eigenvalues and singular values of a given matrix are computed;
- to find out how the largest (and smallest) few eigenvalues and singular values are computed;
- to see some interesting applications of eigenvalues and singular values.

Outline

- Algorithms for a few eigenvalues
- Uses of eigenvalues and eigenvectors
- Uses of SVD

Review of basic concepts

- ▶ Let $A \in \mathbb{R}^{n \times n}$. A scalar λ and a vector $\mathbf{x} \neq \mathbf{0}$ are an **eigenvalue-eigenvector** pair (or eigenpair) if

$$A\mathbf{x} = \lambda\mathbf{x}.$$

- ▶ The set $\sigma(A)$ of all eigenvalues form the **spectrum** of A .
- ▶ The nullspace $\mathcal{N}(\lambda I - A)$ is called **eigenspace** of λ .
 $\dim(\mathcal{N}(\lambda I - A)) = \text{geometric} \leq \text{algebraic multiplicity of } \lambda$.
- ▶ $A\mathbf{x} = \lambda\mathbf{x} \implies (A + \alpha I)\mathbf{x} = (\lambda + \alpha)\mathbf{x}$
- ▶ $A\mathbf{x} = \lambda\mathbf{x} \implies A^k\mathbf{x} = \lambda^k\mathbf{x}$
- ▶ **Similarity transformation**: Given a nonsingular matrix S , the matrix $S^{-1}AS$ has the same eigenvalues as A .
(What about the eigenvectors?)

Review of basic concepts (cont.)

- **Spectral decomposition:** For a diagonalizable $n \times n$ real matrix A there are n (generally complex-valued) eigenpairs $(\lambda_j, \mathbf{x}_j)$, with $X = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n]$ nonsingular.

Then,

$$AX = X\Lambda \iff A = X\Lambda X^{-1}$$

with $\Lambda = \text{diag}(\lambda_1, \lambda_2, \dots, \lambda_n)$.

- Any vector can be written as $\mathbf{y} = X\boldsymbol{\alpha} = \sum_j \mathbf{x}_j \alpha_j$ and

$$A\mathbf{y} = \sum_{j=1}^n \alpha_j A\mathbf{x}_j = \sum_{j=1}^n \alpha_j \lambda_j \mathbf{x}_j.$$

Review of basic concepts (cont.)

- *Does every real $n \times n$ matrix have n eigenvalues?*

We can write the eigenvalue problem as a homogeneous linear system

$$(\lambda I - A)\mathbf{x} = \mathbf{0}.$$

Since we want a nontrivial \mathbf{x} , this means that $\lambda I - A$ must be singular.

Therefore we can find λ by forming the characteristic polynomial and finding its roots.

That is, in principle we solve

$$\det(\lambda I - A) = 0,$$

which has n in general complex, not necessarily distinct roots.

Numerical example

- Example: The matrix

$$A = \begin{pmatrix} -0.9880 & 1.8000 & -0.8793 & -0.5977 & -0.7819 \\ -1.9417 & -0.5835 & -0.1846 & -0.7250 & 1.0422 \\ 0.6003 & -0.0287 & -0.5446 & -2.0667 & -0.3961 \\ 0.8222 & 1.4453 & 1.3369 & -0.6069 & 0.8043 \\ -0.4187 & -0.2939 & 1.4814 & -0.2119 & -1.2771 \end{pmatrix}$$

has eigenvalues given approximately by $\lambda_1 = -2$, $\lambda_2 = -1 + 2.5i$, $\lambda_3 = -1 - 2.5i$, $\lambda_4 = 2i$, and $\lambda_5 = -2i$.

It is known that closed form formulas for the roots of a polynomial do not generally exist if the polynomial is of degree 5 or higher. Thus we cannot expect to be able to solve the eigenvalue problem in a finite procedure.

Singular value decomposition

Let A be real $m \times n$ (rectangular in general). Then there are orthogonal matrices U , V such that

$$A = U\Sigma V^T,$$

where

$$\Sigma = \begin{pmatrix} S & 0 \\ 0 & 0 \end{pmatrix}, \quad S = \text{diag}\{\sigma_1, \dots, \sigma_r\},$$

with the **singular values** $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$, $\sigma_{r+1} = \dots = \sigma_n = 0$.

Connection to eigenvalues: $\sigma_i = \sqrt{\lambda_i}$, where λ_i are eigenvalues of $A^T A$.

Recall eigenvalues and singular values (Ch. 4)

- For a real, square $n \times n$ matrix A , an eigenvalue λ and corresponding eigenvector $\mathbf{x} \neq \mathbf{0}$ satisfy $A\mathbf{x} = \lambda\mathbf{x}$.
- There are n (possibly complex) eigenpairs $\lambda_1, \lambda_2, \dots, \lambda_n$ and eigenvectors $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$ s.t. $A\mathbf{x}_i = \lambda_i\mathbf{x}_i$. For $\Lambda = \text{diag}(\lambda_i)$

$$AX = X\Lambda.$$

- If A is non-defective then the eigenvector matrix X is nonsingular:

$$X^{-1}AX = \Lambda.$$

- For a real, $m \times n$ matrix A , the singular value decomposition (SVD) is

$$A = U\Sigma V^T$$

where U $m \times m$ and V $n \times n$ are orthogonal matrices and Σ is “diagonal” consisting of zeros and singular values $\sigma_1 \geq \sigma_2, \dots \geq \sigma_r > 0$, $r \leq \min(m, n)$.

- Note: singular values are square roots of eigenvalues of $A^T A$.

Singular value decomposition

Theorem (Singular value decomposition (SVD))

For any matrix $A \in \mathbb{R}^{m \times n}$, $m \geq n$, there are unitary $U \in \mathbb{R}^{m \times m}$, $V \in \mathbb{R}^{n \times n}$ and a (generalized) diagonal matrix $\Sigma = (\sigma_1, \dots, \sigma_n)$, $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$, $\sigma_{r+1} = \dots = \sigma_m = 0$, such that

$$A = U \Sigma V^T$$

$$\begin{pmatrix} A \end{pmatrix} = \begin{pmatrix} U \end{pmatrix} \begin{pmatrix} \Sigma \end{pmatrix} \begin{pmatrix} V^T \end{pmatrix}$$

Singular value decomposition (cont.)

The decomposition

$$A = U\Sigma V^T$$

is called **singular value decomposition (SVD)**. There is an economical variant:

$$\begin{pmatrix} \boxed{A} \end{pmatrix} = \begin{pmatrix} \boxed{U} \end{pmatrix} \begin{pmatrix} \boxed{\begin{array}{c} \text{---} \diagup \Sigma \diagdown \text{---} \end{array}} \end{pmatrix} \begin{pmatrix} \boxed{V^T} \end{pmatrix}$$

Singular value decomposition (cont.)

SVD useful to compute

- Condition number of a matrix:

$$\kappa(A) = \kappa_2(A) = \sigma_1(A)/\sigma_p(A).$$

If $\sigma_p(A) = 0$ then the matrix is singular or rank deficient.

- Nullspace of a matrix:

The columns of V corresponding to zero singular values span $\mathcal{N}(A)$.

- Range of a matrix:

The columns of U corresponding to positive singular values span $\mathcal{R}(A)$.

Singular value decomposition (cont.)

MATLAB-functions for computing the SVD:

- `s = svd(A)` : computes singular values of matrix A
- `[U,S,V] = svd(A)` : computes singular value decomposition
- `[U,S,V] = svd(A,0)` : “economical” singular value decomposition
for $m > n$: $U \in \mathbb{R}^{m \times n}$, $\Sigma \in \mathbb{R}^{n \times n}$, $V \in \mathbb{R}^{n \times n}$
- `s = svds(A,k)` : k largest singular values of sparse A
- `[U,S,V] = svds(A,k)` : partial singular value decomposition: $U \in \mathbb{R}^{m \times k}$, $V \in \mathbb{R}^{n \times k}$, $\Sigma \in \mathbb{R}^{k,k}$ diagonal with k largest singular values of A .

MATLAB code for computing nullspace

```
function V = kerncomp(A,tol)
% computes an orthonormal basis of nullspace(A) using SVD.
% kernel selection with relative tolerance tol
if (nargin < 2), tol = eps; end
[U,S,V] = svd(A); % Singular Value Decomposition
s = diag(S);      % Extract vector of singular values
% find singular values of relative size < tol*sigma(1)
V = V(:,find(s < tol*s(1))); % rightmost columns of V.
```

Condition number of rectangular matrices

For a rectangular matrix $A \in \mathbb{R}^{m \times n}$, $m \geq n$, let $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n > 0$ be the eigenvalues of $B = A^T A$. Define **singular values** $\sigma_1, \dots, \sigma_n$ of A as

$$\sigma_i = \sqrt{\lambda_i}, \quad i = 1, 2, \dots, n.$$

Define **condition number** of the rectangular A

$$\kappa(A) = \kappa_2(A) = \frac{\sigma_1}{\sigma_n} = \sqrt{\frac{\lambda_1}{\lambda_n}}.$$

Note

$$\kappa_2(B) = \frac{\lambda_1}{\lambda_n} = \kappa_2(A)^2$$

**B: symmetric
and positive
definite matrix
Why?**

Outline

- Algorithm to find a few eigenvalues: the most dominant eigenvalue
- Uses of eigenvalues and eigenvectors
- Uses of SVD

Power method

- Given \mathbf{v}_0 , expand using eigenpairs

$$\mathbf{v}_0 = \sum_{j=1}^n \beta_j \mathbf{x}_j,$$

where $A\mathbf{x}_j = \lambda_j \mathbf{x}_j$, $j = 1, \dots, n$.

- Then

$$A\mathbf{v}_0 = \sum_{j=1}^n \beta_j A\mathbf{x}_j = \sum_{j=1}^n (\beta_j \lambda_j) \mathbf{x}_j.$$

- Multiplying by A $k-1$ times gives

$$A^k \mathbf{v}_0 = \sum_{j=1}^n \beta_j \lambda_j^{k-1} A\mathbf{x}_j = \sum_{j=1}^n (\beta_j \lambda_j^k) \mathbf{x}_j.$$

Power method (cont.)

Assuming:

- A non-defective
- $|\lambda_1| > |\lambda_j|, j = 2, \dots, n$
- $\beta_1 \neq 0$

Obtain

$$A^k \mathbf{v}_0 \rightarrow \mathbf{x}_1.$$

Given eigenvector, obtain eigenvalue from Rayleigh quotient

$$\lambda_1 = \frac{\mathbf{x}_1^T A \mathbf{x}_1}{\mathbf{x}_1^T \mathbf{x}_1}.$$

Defective matrix: Matrices whose eigenvectors do not span \mathbb{R}^n

Algorithm: Power method

```
Input matrix  $A$  and initial guess  $\mathbf{v}_0$ .  
for  $k := 1, 2, \dots$  until termination do  
     $\tilde{\mathbf{v}} = A\mathbf{v}_{k-1}$   
     $\mathbf{v}_k = \tilde{\mathbf{v}} / \|\tilde{\mathbf{v}}\|$   
     $\lambda_1^{(k)} = \mathbf{v}_k^T A \mathbf{v}_k$   
end for
```

The convergence criterion may be, e.g., the angle between \mathbf{v}_k and \mathbf{v}_{k-1} :

$$\sin \angle(\mathbf{v}_k, \mathbf{v}_{k-1}) = \|\mathbf{v}_k - \mathbf{v}_{k-1}(\mathbf{v}_{k-1}^T \mathbf{v}_k)\| = \|(I - \mathbf{v}_{k-1}\mathbf{v}_{k-1}^T)\mathbf{v}_k\|.$$

Power method algorithm and properties

- Algorithm:

For $k = 1, 2, \dots$ until termination

$$\tilde{\mathbf{v}} = A\mathbf{v}_{k-1}$$

$$\mathbf{v}_k = \tilde{\mathbf{v}} / \|\tilde{\mathbf{v}}\|_2$$

$$\lambda_1^{(k)} = \mathbf{v}_k^T A \mathbf{v}_k.$$

- Properties:


If there is no single dominant eigenvalue, or
eigenvectors are not linearly independent

- Simple, basic, can be slow.
- Can be applied to large, sparse matrices.
- Used as a building block for other, more robust algorithms.

Example

$$A = \text{diag}(1, 2, \dots, 30, 31, 32) \in \mathbb{R}^{32 \times 32}$$

$$B = \text{diag}(1, 2, \dots, 30, 30, 32) \in \mathbb{R}^{32 \times 32}$$



Absolute error:

$$|\lambda_1^{(k)} - \lambda_1| \equiv |\lambda_1^{(k)} - 32|.$$

$$\frac{31}{32} \approx 0.968$$

$$\frac{30}{32} \approx 0.938$$

Truncated SVD

Solving almost singular linear systems

- ▶ Let $A \in \mathbb{R}^{n \times n}$. If $\kappa(A)$ is very large, then solving $A\mathbf{x} = \mathbf{b}$ can be an **ill-conditioned** problem.
- ▶ With the SVD of A we get

$$\mathbf{x} = V\Sigma^{-1}U^T\mathbf{b}.$$

- ▶ If the problem is too ill-conditioned to be of use, then people often seek to **regularize** it. This means, replace the given problem intelligently by a nearby problem which is better conditioned.
- ▶ Using SVD this can be done by setting the singular values below a cutoff tolerance to 0, and minimizing the 2-norm of the solution to the resulting underdetermined problem.

Solving almost singular linear systems (cont.)

► We proceed as follows:

1. Starting from n go backward until r is found such that $\frac{\sigma_1}{\sigma_r}$ is tolerable in size. This is the condition number of the problem that we actually solve.
2. Calculate $\mathbf{z} = U^T \mathbf{b}$; in fact just the first r components of \mathbf{z} are needed. In other words, if $\mathbf{u}_i = U \mathbf{e}_i$ is the i th column vector of U , then $z_i = \mathbf{u}_i^T \mathbf{b}$, $i = 1, \dots, r$.
3. Calculate $y_i = z_i / \sigma_i$, $i = 1, \dots, r$, and set $y_i = 0$, $r < i \leq n$.
4. Calculate $\mathbf{x} = V \mathbf{y}$. This really involves only the first r columns of V and the first r components of \mathbf{y} . In other words, if \mathbf{v}_i is the i th column vector of V , then $\mathbf{x} = \sum_{i=1}^r y_i \mathbf{v}_i$.

Note: In general, the resulting \mathbf{x} does not satisfy $A\mathbf{x} = \mathbf{b}$. But it's the best one can do under certain circumstances, and it produces a solution \mathbf{x} of the smallest norm for a sufficiently well-conditioned approximate problem.

Truncated SVD and data compression

- Given an $m \times n$ matrix A , the best rank- r approximation of $A = U\Sigma V^T$ is the matrix

$$A_r = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^T.$$

- This is another example of truncated SVD (TSVD), so named because only the first r columns of U and V are utilized.
- It is a **model reduction** technique, which is a best approximation in the sense that $\|A - A_r\|_2$ is minimal over all possible rank- r matrices. The minimum residual norm is equal to σ_{r+1} .
- Note A_r uses only $r(m + n + 1)$ storage locations – significantly fewer than mn if $r \ll \min(m, n)$.

Example

- The least squares problem $\min_{\mathbf{x}} \|C\mathbf{x} - \mathbf{b}\|$ is easily solved for

$$C = \begin{pmatrix} 1 & 0 & 1 \\ 2 & 3 & 5 \\ 5 & 3 & -2 \\ 3 & 5 & 4 \\ -1 & 6 & 3 \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} 4 \\ -2 \\ 5 \\ -2 \\ 1 \end{pmatrix}.$$

Call the solution $\hat{\mathbf{x}}$.

- Next consider $\min_{\mathbf{x}} \|A\mathbf{x} - \mathbf{b}\|$, where we add to C a column that is sum of the three previous ones:

$$A = \begin{pmatrix} 1 & 0 & 1 & 2 \\ 2 & 3 & 5 & 10 \\ 5 & 3 & -2 & 6 \\ 3 & 5 & 4 & 12 \\ -1 & 6 & 3 & 8 \end{pmatrix}.$$

- Note that A has $m = 5$, $n = 4$, $r = 3$: can't reliably solve this using normal equations or even QR.
- But the truncated SVD method works: if \mathbf{x} solves the problem with A , obtain $\|\mathbf{x}\| \approx \|\hat{\mathbf{x}}\|$, and $\|A\mathbf{x} - \mathbf{b}\| \approx \|C\hat{\mathbf{x}} - \mathbf{b}\|$.

Example

- The least squares problem $\min_{\mathbf{x}} \|C\mathbf{x} - \mathbf{b}\|$ is easily solved for

$$C = \begin{pmatrix} 1 & 0 & 1 \\ 2 & 3 & 5 \\ 5 & 3 & -2 \\ 3 & 5 & 4 \\ -1 & 6 & 3 \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} 4 \\ -2 \\ 5 \\ -2 \\ 1 \end{pmatrix}.$$

Call the solution $\hat{\mathbf{x}}$.

- Next consider $\min_{\mathbf{x}} \|A\mathbf{x} - \mathbf{b}\|$, where we add to C a column that is sum of the three previous ones:

$$A = \begin{pmatrix} 1 & 0 & 1 & 2 \\ 2 & 3 & 5 & 10 \\ 5 & 3 & -2 & 6 \\ 3 & 5 & 4 & 12 \\ -1 & 6 & 3 & 8 \end{pmatrix}.$$

- Note that A has $m = 5$, $n = 4$, $r = 3$: can't reliably solve this using normal equations.
- But the truncated SVD method works: if \mathbf{x} solves the problem with A , obtain $\|\mathbf{x}\| \approx \|\hat{\mathbf{x}}\|$, and $\|A\mathbf{x} - \mathbf{b}\| \approx \|C\hat{\mathbf{x}} - \mathbf{b}\|$.

Compressing image information

- ▶ Image composed of $m \times n$ pixels (greyscale, BMP format).
- ▶ Matrix $A \in \mathbb{R}^{m \times n}$, $0 \leq a_{ij} \leq 255$.
- ▶ Compression scheme using SVD: Replace A by A_k . $k = ?$
- ▶ Represent A_k by storing the first k columns of U and V .
- ▶ Instead of mn numbers only store $k(m + n)$ numbers.


Compressing image information (cont.)

```
close all
colormap('gray')
load clown.mat;
figure(1)
image(X);

[U,S,V] = svd(X);
figure(2)
r = 20;
colormap('gray')
image(U(:,1:r)*S(1:r,1:r)*V(:,1:r)')
```


Example

Consider the following clown image, taken from MATLAB's image repository:



Compressed image of clown example

Here is what we get if we use $r = 20$.



Assessment

- The original image requires $200 \times 320 = 64000$ matrix entries to be stored; the compressed image requires merely $20 \cdot (200 + 320 + 1) \approx 10000$ storage locations.
- By storing less than 16% of the data, we get a reasonable image. It's not great, but all main the features of the clown are clearly shown.
- What are less clear in the compressed image are fine features (high frequency), such as the fine details of the clown's hair.
- Certainly, [advanced](#) techniques such as [DCT](#) (Chapter 13) and [wavelet](#) are far superior to SVD for the task of image compression. Still, our example visually shows that most information is stored already in the leading singular vectors.